

part_2_orange

January 23, 2025

0.0.1 Introduction

In this part of the analysis, we build upon the clustering results obtained in the first step, where we identified groups using the best-performing clustering method for this dataset. We then compared two algorithms, Random Forest and Gradient Boosting, to develop an effective classification algorithm capable of assigning individuals to the identified groups based on their responses. Random Forest was chosen for its robustness, interpretability, and ability to handle diverse data types, while Gradient Boosting was selected for its strong predictive performance and capacity to capture non-linear relationships. Finally, we optimized the number of questions used in the classification process by analyzing feature importance and iteratively reducing the feature set to achieve the best trade-off between performance and simplicity. This step ensures that the resulting algorithm remains both efficient and accurate, providing a practical solution for future use with minimal inputs.

0.1 IMPORT

```
[12]: import matplotlib.pyplot as plt
      from sklearn.metrics import accuracy_score
      import pandas as pd
      from sklearn.ensemble import GradientBoostingClassifier
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.preprocessing import StandardScaler
      from sklearn.preprocessing import StandardScaler
      from sklearn.cluster import KMeans
```

0.2 Data Loading

```
[13]: file_path = "fic_epita_kantar_codes.csv"
      data_codes = pd.read_csv(file_path, sep=';')
```

Select relevant columns

```
[14]: columns_of_interest = [
      "A9_1_slice", "A9_2_slice", "A9_3_slice", "A9_4_slice", "A9_5_slice",
      "A9_6_slice", "A9_7_slice", "A9_8_slice", "A9_9_slice", "A9_10_slice",
      "A9_11_slice", "A9_12_slice", "A9_13_slice", "A9_14_slice", "A9_15_slice",
      "A9_16_slice", "A10_1_slice", "A10_2_slice", "A10_3_slice", "A10_4_slice",
```

```

    "A10_5_slice", "A10_6_slice", "A10_7_slice", "A10_8_slice",
    "A11_1_slice", "A11_2_slice", "A11_3_slice", "A11_4_slice", "A11_5_slice",
    "A11_6_slice", "A11_7_slice", "A11_8_slice", "A11_9_slice", "A11_10_slice",
    "A11_11_slice", "A11_12_slice", "A11_13_slice"
]
subset_data = data_codes[columns_of_interest]

```

0.3 Adding the cluster column using the best algo with the best number of cluster found in part 1

```

[ ]: # Scale the data
scaler = StandardScaler()
scaled_data = scaler.fit_transform(subset_data.fillna(0)) # Replace missing
↳ values

# Perform KMeans clustering
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans_labels = kmeans.fit_predict(scaled_data)

# Add cluster labels to the data
subset_data['cluster'] = kmeans_labels

```

/tmp/ipykernel_131606/3220641698.py:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
subset_data['cluster'] = kmeans_labels

0.4 Training Random Forest

```

[16]: # Define features (golden questions) and target (clusters)
X = subset_data.drop(columns=['cluster'])
y = subset_data['cluster']

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Train and Evaluate model
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
print("Classification Report:\n", classification_report(y_test, y_pred))

```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.93	0.97	0.95	262
1	0.98	0.90	0.94	184
2	0.95	0.96	0.96	554
accuracy			0.95	1000
macro avg	0.96	0.94	0.95	1000
weighted avg	0.95	0.95	0.95	1000

0.5 Training Gradient Boosting

```
[17]: # Define features (questions) and target (clusters)
X2 = subset_data.drop(columns=['cluster'])
y2 = subset_data['cluster']

# Split into training and test sets
X_train2, X_test2, y_train2, y_test2 = train_test_split(X2, y2, test_size=0.2,
    random_state=42)

# Train and Evaluate model
gb = GradientBoostingClassifier(random_state=42)
gb.fit(X_train2, y_train2)
y_pred2 = gb.predict(X_test2)
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.97	0.95	262
1	0.98	0.90	0.94	184
2	0.95	0.96	0.96	554
accuracy			0.95	1000
macro avg	0.96	0.94	0.95	1000
weighted avg	0.95	0.95	0.95	1000

0.6 Getting the importance of each question

For random forest

```
[18]: feature_importance = rf.feature_importances_
importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': feature_importance
}).sort_values(by='Importance', ascending=False)
```

```
print("Feature Importance (Top Questions) for Random Forest:\n", importance_df)
```

Feature Importance (Top Questions) for Random Forest:

	Feature	Importance
33	A11_10_slice	0.083317
30	A11_7_slice	0.071852
28	A11_5_slice	0.071477
26	A11_3_slice	0.056607
31	A11_8_slice	0.052130
7	A9_8_slice	0.043957
13	A9_14_slice	0.040699
36	A11_13_slice	0.039214
20	A10_5_slice	0.038229
3	A9_4_slice	0.035639
34	A11_11_slice	0.035226
2	A9_3_slice	0.029591
1	A9_2_slice	0.029545
0	A9_1_slice	0.026178
24	A11_1_slice	0.023914
25	A11_2_slice	0.022092
9	A9_10_slice	0.021260
32	A11_9_slice	0.020025
35	A11_12_slice	0.019942
27	A11_4_slice	0.019370
16	A10_1_slice	0.018899
6	A9_7_slice	0.018642
10	A9_11_slice	0.017545
15	A9_16_slice	0.016899
12	A9_13_slice	0.016795
11	A9_12_slice	0.016338
14	A9_15_slice	0.016270
22	A10_7_slice	0.014869
23	A10_8_slice	0.013685
8	A9_9_slice	0.011098
29	A11_6_slice	0.009880
19	A10_4_slice	0.009809
17	A10_2_slice	0.008753
5	A9_6_slice	0.008104
21	A10_6_slice	0.007870
18	A10_3_slice	0.007530
4	A9_5_slice	0.006751

For Gradient Boosting

```
[19]: feature_importance2 = gb.feature_importances_  
importance_df2 = pd.DataFrame({  
    'Feature': X2.columns,  
    'Importance': feature_importance2
```

```
}).sort_values(by='Importance', ascending=False)

print("Feature Importance (Top Questions) for Gradient Boosting:\n",
      importance_df2)
```

Feature Importance (Top Questions) for Gradient Boosting:

	Feature	Importance
28	A11_5_slice	0.140620
30	A11_7_slice	0.112928
33	A11_10_slice	0.098550
7	A9_8_slice	0.064707
26	A11_3_slice	0.060625
31	A11_8_slice	0.055011
13	A9_14_slice	0.048094
3	A9_4_slice	0.043571
34	A11_11_slice	0.042199
36	A11_13_slice	0.037003
20	A10_5_slice	0.036868
0	A9_1_slice	0.031365
2	A9_3_slice	0.030072
1	A9_2_slice	0.028438
16	A10_1_slice	0.023673
35	A11_12_slice	0.017286
27	A11_4_slice	0.015426
9	A9_10_slice	0.014880
24	A11_1_slice	0.014104
6	A9_7_slice	0.013802
10	A9_11_slice	0.012057
25	A11_2_slice	0.011214
12	A9_13_slice	0.008521
32	A11_9_slice	0.007891
15	A9_16_slice	0.006086
22	A10_7_slice	0.005294
14	A9_15_slice	0.005157
23	A10_8_slice	0.002945
19	A10_4_slice	0.002839
17	A10_2_slice	0.002215
11	A9_12_slice	0.001998
8	A9_9_slice	0.001505
29	A11_6_slice	0.001439
5	A9_6_slice	0.000662
21	A10_6_slice	0.000586
18	A10_3_slice	0.000263
4	A9_5_slice	0.000106

0.7 Finding performance when lowering question amount

For random forest

```
[20]: results = []
      for i in range(1, len(importance_df) + 1):
          top_features = importance_df['Feature'].head(i)
          X_train_reduced = X_train[top_features]
          X_test_reduced = X_test[top_features]

          rf_reduced = RandomForestClassifier(random_state=42)
          rf_reduced.fit(X_train_reduced, y_train)
          y_pred_reduced = rf_reduced.predict(X_test_reduced)

          acc = accuracy_score(y_test, y_pred_reduced)
          results.append({'Num_Features': i, 'Accuracy': acc})

      results_df = pd.DataFrame(results)
      print(results_df)
```

	Num_Features	Accuracy
0	1	0.730
1	2	0.749
2	3	0.795
3	4	0.800
4	5	0.806
5	6	0.830
6	7	0.836
7	8	0.836
8	9	0.841
9	10	0.875
10	11	0.881
11	12	0.882
12	13	0.889
13	14	0.893
14	15	0.903
15	16	0.901
16	17	0.906
17	18	0.902
18	19	0.906
19	20	0.919
20	21	0.925
21	22	0.929
22	23	0.937
23	24	0.939
24	25	0.942
25	26	0.942
26	27	0.944
27	28	0.940
28	29	0.947
29	30	0.950
30	31	0.948

31	32	0.948
32	33	0.956
33	34	0.955
34	35	0.957
35	36	0.951
36	37	0.948

For Gradient Boosting

```
[21]: results2 = []
      for i in range(1, len(importance_df2) + 1):
          top_features2 = importance_df2['Feature'].head(i)
          X_train_reduced2 = X_train2[top_features2]
          X_test_reduced2 = X_test2[top_features2]

          gb_reduced = GradientBoostingClassifier(random_state=42)
          gb_reduced.fit(X_train_reduced2, y_train2)
          y_pred_reduced2 = gb_reduced.predict(X_test_reduced2)


          acc = accuracy_score(y_test2, y_pred_reduced2)
          results2.append({'Num_Features': i, 'Accuracy': acc})

      results_df2 = pd.DataFrame(results2)
      print(results_df2)
```

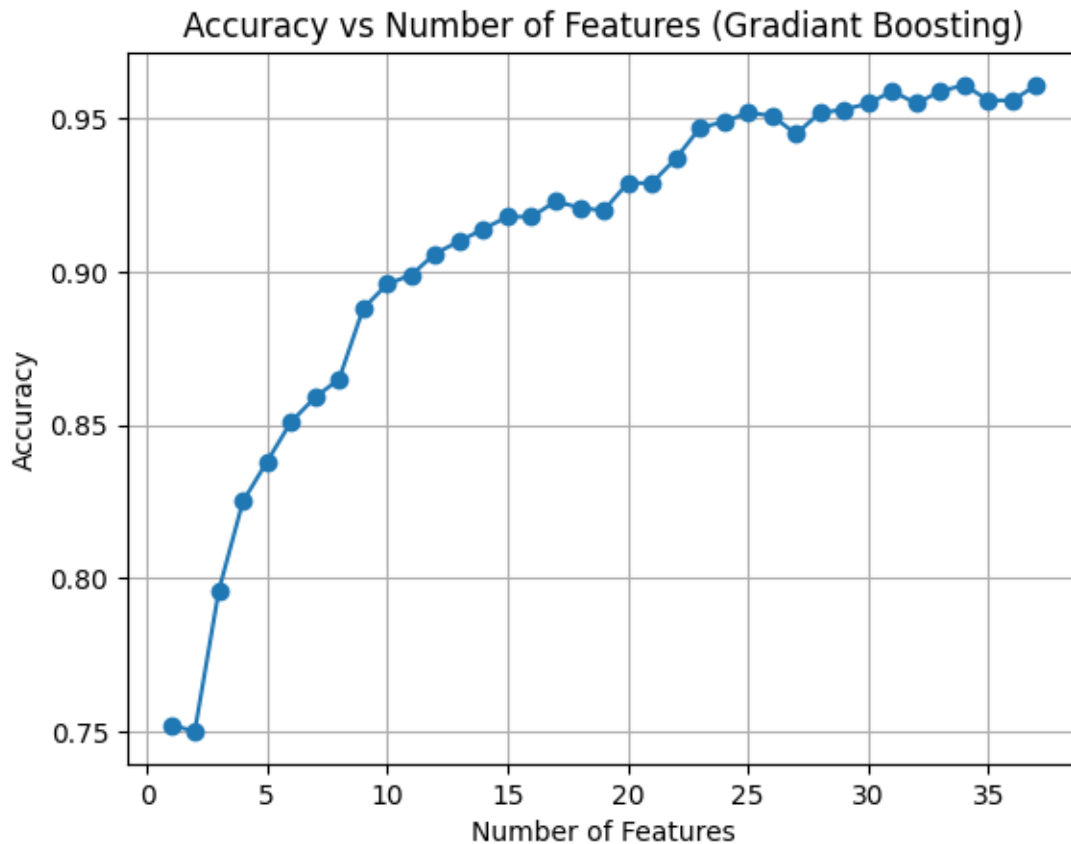
	Num_Features	Accuracy
0	1	0.752
1	2	0.750
2	3	0.796
3	4	0.825
4	5	0.838
5	6	0.851
6	7	0.859
7	8	0.865
8	9	0.888
9	10	0.896
10	11	0.899
11	12	0.906
12	13	0.910
13	14	0.914
14	15	0.918
15	16	0.918
16	17	0.923
17	18	0.921
18	19	0.920
19	20	0.929
20	21	0.929
21	22	0.937
22	23	0.947

23	24	0.949
24	25	0.952
25	26	0.951
26	27	0.945
27	28	0.952
28	29	0.953
29	30	0.955
30	31	0.959
31	32	0.955
32	33	0.959
33	34	0.961
34	35	0.956
35	36	0.956
36	37	0.961

```
[22]: plt.plot(results_df['Num_Features'], results_df['Accuracy'], marker='o')
plt.xlabel("Number of Features")
plt.ylabel("Accuracy")
plt.title("Accuracy vs Number of Features (Random Forest)")
plt.grid()
plt.show()
```




```
[23]: plt.plot(results_df2['Num_Features'], results_df2['Accuracy'], marker='o')
plt.xlabel("Number of Features")
plt.ylabel("Accuracy")
plt.title("Accuracy vs Number of Features (Gradient Boosting)")
plt.grid()
plt.show()
```



0.7.1 Conclusion

Both Random Forest and Gradient Boosting are similarly effective, with Gradient Boosting being slightly better in terms of accuracy. Therefore, the optimal choice is Gradient Boosting.

Using 23 features results in an accuracy nearing 95%, striking the best balance between performance and complexity. Beyond that, the improvements in accuracy become minimal, and the model starts to slow down, making it the optimal choice for maximizing classification performance while maintaining efficiency.